

Account Book Triage

Turns a stale CRM account export into a ranked list of what a rep should do next — and a separate list of records nobody should touch until a human resolves them.

Grounded in Apollo.io company data 7 live endpoints 32 offline tests 0 credits in CI

THE PROBLEM

Reps work their book alphabetically because nothing tells them not to

A rep opens 200 accounts on Monday. A third are missing firmographics. Some were last verified two years ago. A handful point at companies that have since been acquired. Nothing in the CRM separates the account hiring six Account Executives this month from the one dormant since 2023.

So they start at the top of the list, and ops fields another “the data is bad” complaint.

Most prioritization tools answer *which accounts are hot* and quietly assume the records are trustworthy. This one answers two questions and refuses to blend them into a single number:

Is this record fit to act on?

data_health 0–100 — field completeness, days since last verified, and whether the CRM agrees with the source of truth.

Is something happening here?

opportunity 0–100 — GTM hiring intensity, posting recency, news recency, ICP fit.

A hot account with a broken record isn't a medium-priority account.
It's a hot account you can't route.

WHAT IT PRODUCES


Four actions, and one of them is a stop sign

#	ACCOUNT	ACTION	OPP	HEALTH	NEXT BEST ACTION
1	HubSpot	ENGAGE_NOW	85	63	15 of 25 open roles are GTM. Newest posting today.
4	Gong	ENGAGE_NOW	78	93	Hiring Mid-Market AE, France. Posting 1 day old.
9	Drift	FIX_DATA	72	50	Do not route — identity conflict on domain, account_name.
14	Clari	NURTURE	40	100	One open role, 462 days old. Clean record, cold company.
15	Segment	DEPRIORITIZE	32	62	No live hiring or news signal this quarter.

Real output, abridged from a 15-account book. Row 9 is the point. Drift scores 72 on opportunity and is still blocked, because Apollo resolves drift.com to Salesloft — Drift was acquired and the CRM never absorbed it. A rep working that row calls the wrong company. Score alone would have sent them.

HOW IT WORKS

Scripts do the arithmetic. The skill carries the judgment.



Every weight lives in `assets/scoring.yaml` as an annotated config a RevOps lead can open, disagree with, edit, and re-run. The model never computes a score, never invents a weight, and never phrases a next-best-action — those are templates fed by structured data, because the inputs are already structured and an LLM would only add a hallucination surface.

Why a skill rather than a script

The deterministic work belongs in scripts, and it's there. What a bare script can't carry is the procedure around it: dry-run before spending, read the credit cost aloud before a live run, never auto-apply a conflict, and distinguish three kinds of domain mismatch that have opposite remedies. That's knowledge an operator needs at the moment of use.

Why not point an agent at the Apollo MCP directly

Both paths were tested, so this is an observation rather than an argument. MCP gives an agent tool access, not a method. Asked to “clean up my accounts” against raw tools, a model enriches and then reasons its way to a ranking in prose — with weights nobody can audit and an order that changes between runs. A rubric in a file is reproducible, arguable, and survives the model being swapped underneath it.

It also buys no access. People search through the Apollo MCP returns the identical Free-plan refusal as the REST endpoint — `API_INACCESSIBLE`. **MCP is a transport, not a permission:** the plan gates at the API layer however a client reaches it.

Gaps and conflicts are not the same thing

Filling an empty field from a vendor is safe. Overwriting a populated field because a vendor disagrees is how you destroy a CRM — the rep who typed “Bozeman” may have been on a call last week, and the vendor may be wrong. Gaps auto-patch; conflicts go to a human with both values and a reason.

WHERE AI WAS WRONG

Four things it got confidently wrong, and how they surfaced

AI wrote most of this code. These are the parts worth talking about.

1 It believed the documentation over the API

Apollo’s docs say free accounts need a work-email address for search and enrichment. Reasoning from that, I concluded a Gmail-registered account was blocked and advised moving domains. One live call disproved it in thirty seconds — enrichment works fine, and the real gate is *plan tier*, stated plainly in the 403 body. The proposed fix would have cost an hour and fixed nothing.

2 The first rubric was useless and looked reasonable

It ranked 13 of 15 accounts `ENGAGE_NOW`. Two defects: hiring-recency gave full marks to almost everyone, because companies this size post roles continuously; and hiring intensity saturated at a 40% GTM mix — roughly the *normal* mix for B2B SaaS, so it measured “is a software company.” Reading the code revealed neither. Only the output distribution did.

3 A test asserted the wrong domain fact

“Revenue Manager” was written in as a positive GTM-title case. It’s usually revenue *recognition* — an accounting role — and it appears in Snowflake’s live postings. Counting it would have inflated their hiring signal with a finance hire. The code was right; the test was wrong.

4 Live data forced a distinction I’d have flattened

Apollo returns a `.vercel.app` preview host as Gong’s primary domain (vendor wrong, CRM right), `outreach.ai` for `outreach.io` (same brand), and Salesloft for `drift.com` (real acquisition). One “domain mismatch” rule would have either blocked two good accounts or let the acquired one through.

The pattern: AI-generated logic fails hardest where it reads most plausibly. Checking the `output distribution` caught what reviewing the code did not.

CONSTRAINTS & ASSUMPTIONS

What the free tier actually allows

Verified with live calls before any design work, which is what kept the build inside the lines:

- Company-level endpoints work — enrichment, bulk enrichment, org search, job postings, news, account write-back.
- Every person-level endpoint returns 403 — people search, person enrichment, email and phone lookup. Paid plans only.
- 75 credits per month, one per company enriched. A 200-account book can't be done in one pass on free, so the client enforces a spend budget and caches for 14 days.
- bulk_enrich caps at 10 domains per call — undocumented; the API said so when a 15-domain payload failed.
- The MCP connector changes none of this. Same endpoints, same refusals. It did grant a one-time 100-credit bonus, which is onboarding rather than entitlement.

That entitlement check is why this is an account-level tool. The exercise's suggested "ranked list of key people to reach out to" is not buildable on a free account — contact mapping is named as the paid-tier extension rather than discovered as a wall halfway through.

PRODUCTION HARDENING

What would have to change before reps depend on it

- Rubric governance. `scoring.yaml` decides who a sales team calls. The weights here are informed guesses; they should be backtested against closed-won history, not reasoned about, and changed under review.
- Auth. The POC uses a master key for speed, deliberately — entitlements were unknown, so scoping up front would have locked in guesses. Production needs a scoped key limited to the seven endpoints actually called, rotated, in a secrets manager.
- Write-back safety. Dry-run by default and gap-fills only today. Needs an idempotency key, a reversible audit log, and a staged rollout before it points at a live CRM.
- Signal quality. Apollo's news feed carries real listicle spam — "Best enterprise web hosting platforms compared" is in the captured sample. Filter by category and source reputation, or replace it with a genuine

intent source.

- **Scale.** In-memory and single-threaded. Fine at 200 accounts, wrong at 50,000.
-

Ryan Collins Senior Business Systems Analyst, GTM Built with Claude Code